

```
In [ ]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
from sklearn.ensemble import GradientBoostingRegressor
```

```
In [ ]: data = pd.read_csv("sickness_table.csv")
```

```
In [ ]: data.head(10)
```

```
Out[ ]:
```

	Unnamed: 0	date	n_sick	calls	n_duty	n_sby	sby_need	dafted
0	0	2016-04-01	73	8154.0	1700	90	4.0	0.0
1	1	2016-04-02	64	8526.0	1700	90	70.0	0.0
2	2	2016-04-03	68	8088.0	1700	90	0.0	0.0
3	3	2016-04-04	71	7044.0	1700	90	0.0	0.0
4	4	2016-04-05	63	7236.0	1700	90	0.0	0.0
5	5	2016-04-06	70	6492.0	1700	90	0.0	0.0
6	6	2016-04-07	64	6204.0	1700	90	0.0	0.0
7	7	2016-04-08	62	7614.0	1700	90	0.0	0.0
8	8	2016-04-09	51	5706.0	1700	90	0.0	0.0
9	9	2016-04-10	54	6606.0	1700	90	0.0	0.0

Removing Unwanted Columns

```
In [ ]: # 1.1 Remove Redundant Columns
# Drop the 'Unnamed: 0' column (if it exists)
if 'Unnamed: 0' in data.columns:
    sickness_data = data.drop(columns=['Unnamed: 0'])
```

```
In [ ]: # Convert the 'date' column to datetime format
sickness_data['date'] = pd.to_datetime(sickness_data['date'])
```

```
In [ ]: sickness_data.head()
```

```
Out[ ]:
```

	date	n_sick	calls	n_duty	n_sby	sby_need	dafted
0	2016-04-01	73	8154.0	1700	90	4.0	0.0
1	2016-04-02	64	8526.0	1700	90	70.0	0.0
2	2016-04-03	68	8088.0	1700	90	0.0	0.0
3	2016-04-04	71	7044.0	1700	90	0.0	0.0
4	2016-04-05	63	7236.0	1700	90	0.0	0.0

```
In [ ]: sickness_data.shape
```

Out[]: (1152, 7)

In []: sickness_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1152 entries, 0 to 1151
Data columns (total 7 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   date        1152 non-null   datetime64[ns]
 1   n_sick       1152 non-null   int64
 2   calls       1152 non-null   float64
 3   n_duty      1152 non-null   int64
 4   n_sby       1152 non-null   int64
 5   sby_need    1152 non-null   float64
 6   dafted      1152 non-null   float64
dtypes: datetime64[ns](1), float64(3), int64(3)
memory usage: 63.1 KB
```

In []: sickness_data.describe()

	date	n_sick	calls	n_duty	n_sby	sby_need	dafted
count	1152	1152.000000	1152.000000	1152.000000	1152.0	1152.000000	1152.000000
mean	2017-10-28 12:00:00	68.808160	7919.531250	1820.572917	90.0	34.718750	16.335938
min	2016-04-01 00:00:00	36.000000	4074.000000	1700.000000	90.0	0.000000	0.000000
25%	2017-01-13 18:00:00	58.000000	6978.000000	1800.000000	90.0	0.000000	0.000000
50%	2017-10-28 12:00:00	68.000000	7932.000000	1800.000000	90.0	0.000000	0.000000
75%	2018-08-12 06:00:00	78.000000	8827.500000	1900.000000	90.0	12.250000	0.000000
max	2019-05-27 00:00:00	119.000000	11850.000000	1900.000000	90.0	555.000000	465.000000
std	NaN	14.293942	1290.063571	80.086953	0.0	79.694251	53.394089

In []: sickness_data.isnull().sum()

```
date      0
n_sick    0
calls     0
n_duty    0
n_sby     0
sby_need  0
dafted    0
dtype: int64
```

Data Exploration

```
In [ ]: # Columns of interest
columns_to_analyze = ['n_sick', 'calls', 'sby_need']

# Histograms for Distribution Analysis
for column in columns_to_analyze:
    plt.figure(figsize=(10, 5))
```

```

sns.histplot(sickness_data[column], kde=True, bins=30)
plt.title(f'Distribution of {column}')
plt.xlabel(column)
plt.ylabel('Frequency')
plt.show()


# Boxplots for Outlier Analysis
for column in columns_to_analyze:
    plt.figure(figsize=(10, 5))
    sns.boxplot(x=sickness_data[column])
    plt.title(f'Boxplot of {column}')
    plt.xlabel(column)
    plt.show()

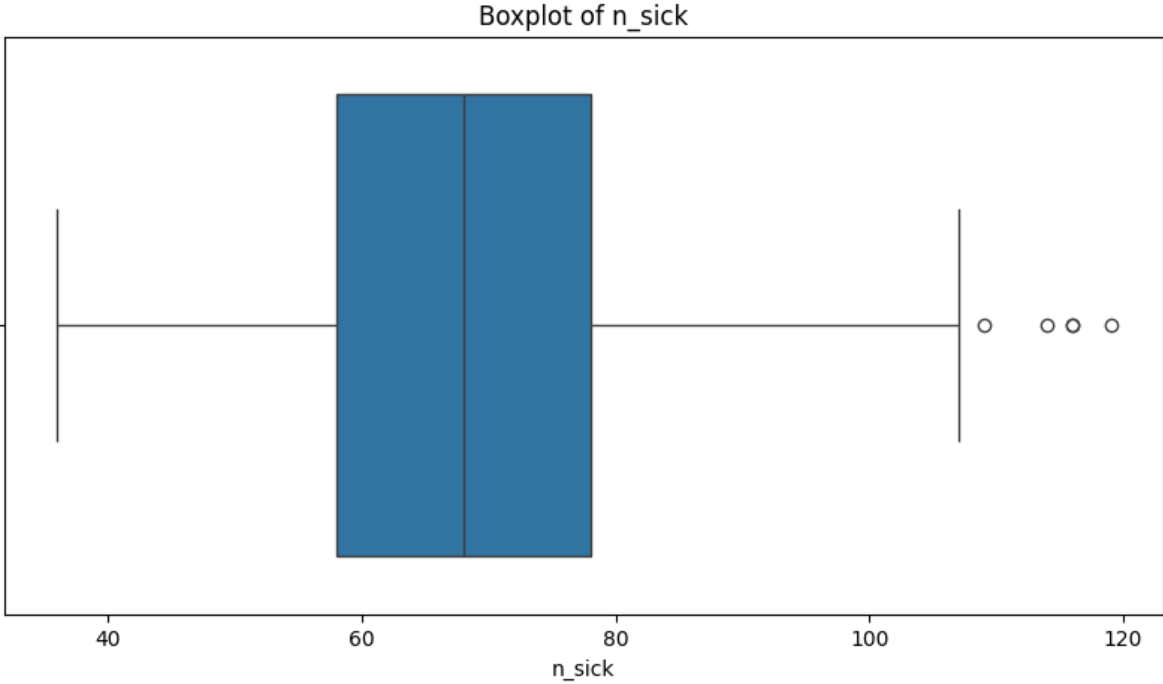
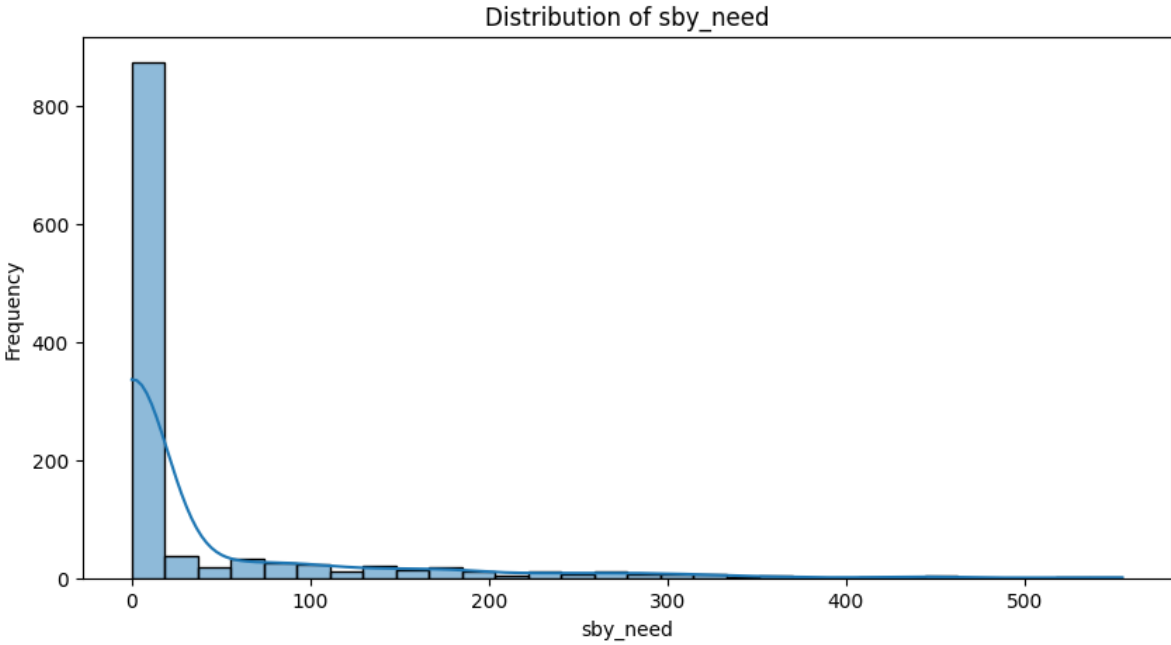

# Summary Statistics
summary_stats = sickness_data[columns_to_analyze].describe().T[['mean', '50%', 'std']]
print(summary_stats)


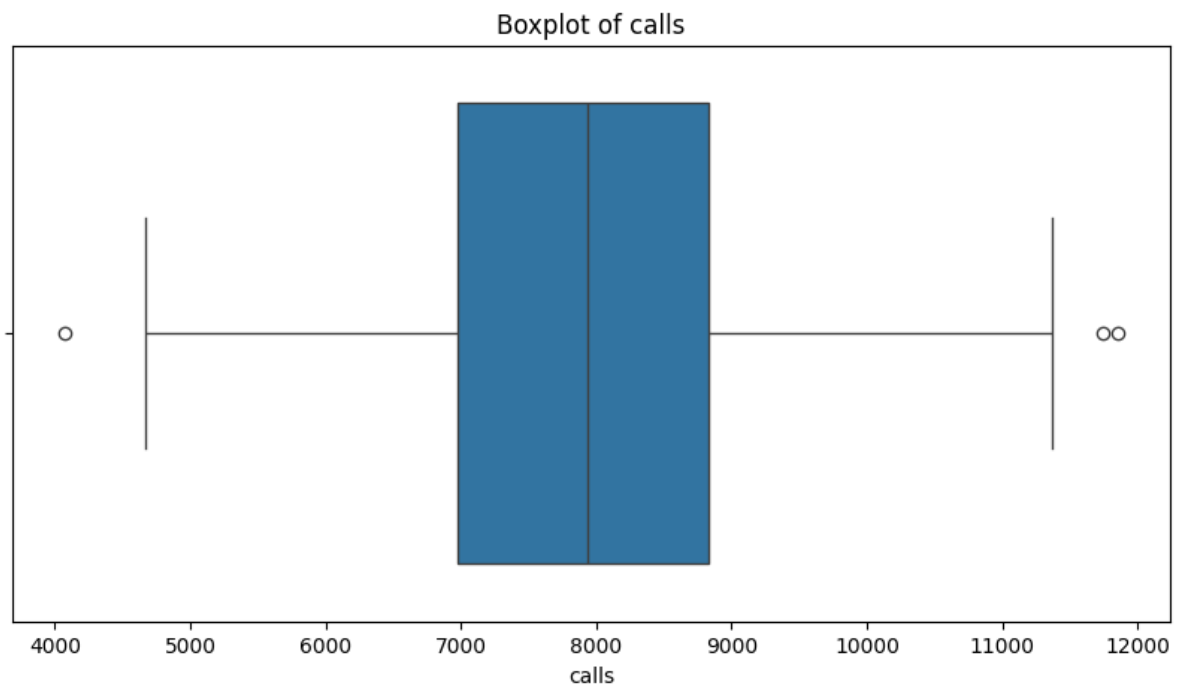
# IQR for Outlier Detection
for column in columns_to_analyze:
    Q1 = sickness_data[column].quantile(0.25)
    Q3 = sickness_data[column].quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - 1.5 * IQR
    upper_bound = Q3 + 1.5 * IQR

    outliers = sickness_data[(sickness_data[column] < lower_bound) | (sickness_data[column] > upper_bound)]
    print(f"Number of outliers detected in {column}: {len(outliers)}")

```








	mean	50%	std
n_sick	68.80816	68.0	14.293942
calls	7919.53125	7932.0	1290.063571
sby_need	34.71875	0.0	79.694251

Number of outliers detected in n_sick: 5
 Number of outliers detected in calls: 3
 Number of outliers detected in sby_need: 256

```
In [ ]: # Time Series Analysis
plt.figure(figsize=(18, 12))

# Plotting n_sick over time
plt.subplot(3, 1, 1)
plt.plot(sickness_data['date'], sickness_data['n_sick'], label='Number of Drivers Sick')
plt.title('Trend of Number of Drivers Sick Over Time')
plt.xlabel('Date')
plt.ylabel('Number of Drivers Sick')
plt.legend()
```


```
plt.tight_layout()
plt.show()
```



```
In [ ]: # Time Series Analysis
plt.figure(figsize=(18, 12))

# Plotting calls over time
plt.subplot(3, 1, 2)
plt.plot(sickness_data['date'], sickness_data['calls'], label='Emergency Calls', color='orange')
plt.title('Trend of Emergency Calls Over Time')
plt.xlabel('Date')
plt.ylabel('Number of Calls')
plt.legend()


plt.tight_layout()
plt.show()
```



```
In [ ]: # Time Series Analysis
plt.figure(figsize=(18, 12))

# Plotting standby drivers activated
plt.subplot(3, 1, 1)
plt.plot(sickness_data['date'], sickness_data['sby_need'], label='Emergency Calls', color='green')
plt.title('Trend of Standby Drivers Activated Over Time')
plt.xlabel('Date')
plt.ylabel('Standby Drivers Activated')
plt.legend()

plt.tight_layout()
plt.show()
```




```
In [ ]: # Time Series Analysis
plt.figure(figsize=(18, 12))

# Plotting drafted over time
```

```
plt.subplot(3, 1, 1)
plt.plot(sickness_data['date'], sickness_data['dafted'], label='Additional Drivers Needed')
plt.title('Trend of Additional Drivers Needed Over Time')
plt.xlabel('Date')
plt.ylabel('Number of Additional Drivers Needed')
plt.legend()

plt.tight_layout()
plt.show()
```



```
In [ ]: sickness_data.corr()
```

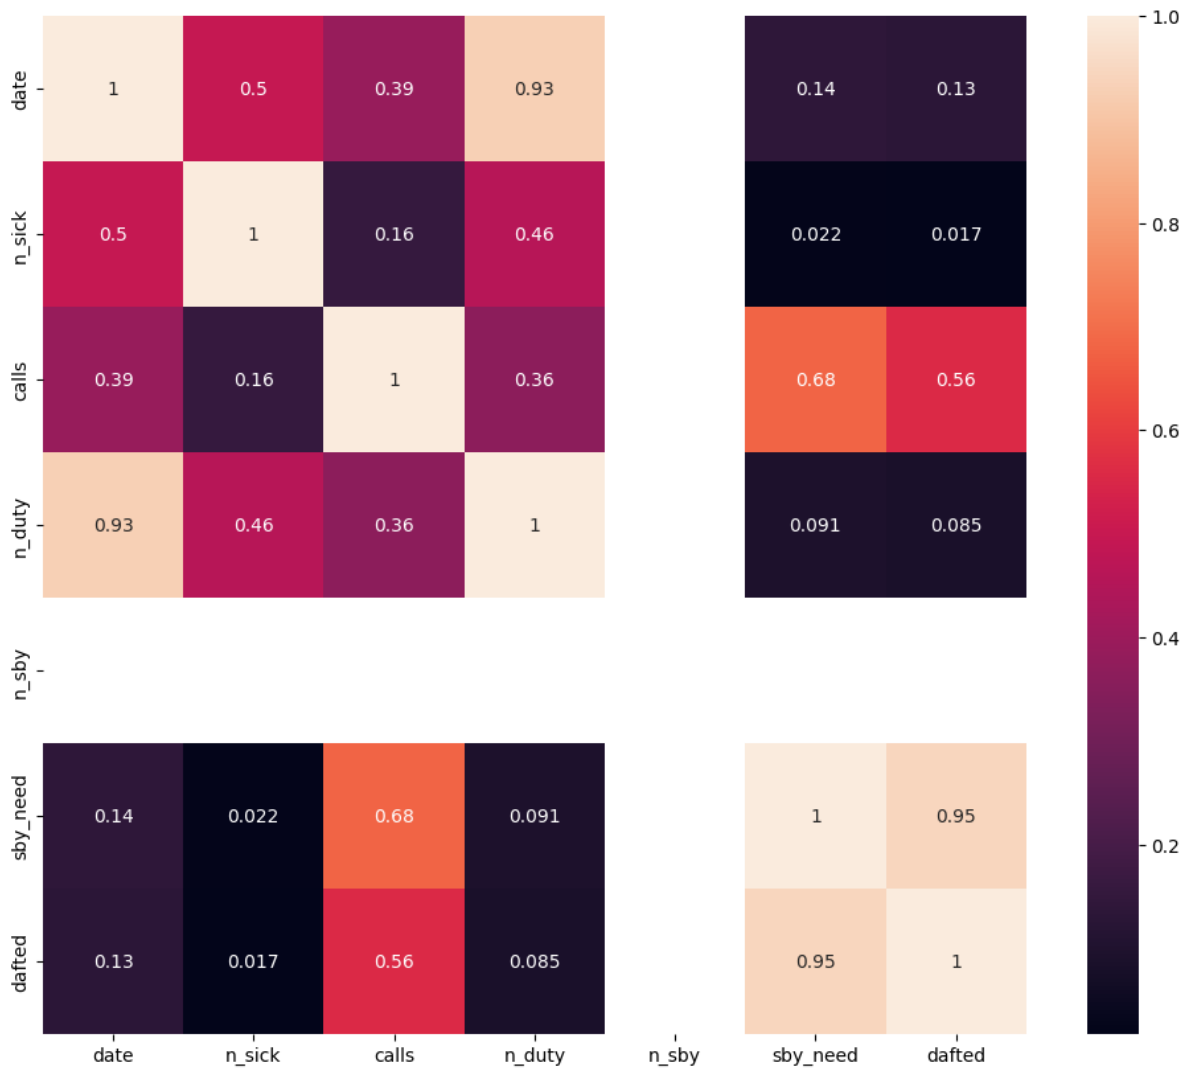
```
Out[ ]:
```

	date	n_sick	calls	n_duty	n_sby	sby_need	dafted
date	1.000000	0.495959	0.385679	0.927437	NaN	0.137543	0.131938
n_sick	0.495959	1.000000	0.155371	0.459501	NaN	0.022321	0.016800
calls	0.385679	0.155371	1.000000	0.364135	NaN	0.677468	0.557340
n_duty	0.927437	0.459501	0.364135	1.000000	NaN	0.090654	0.084955
n_sby	NaN	NaN	NaN	NaN	NaN	NaN	NaN
sby_need	0.137543	0.022321	0.677468	0.090654	NaN	1.000000	0.945168
dafted	0.131938	0.016800	0.557340	0.084955	NaN	0.945168	1.000000

```
In [ ]: plt.figure(figsize= (12, 10))

sns.heatmap(sickness_data.corr(), annot = True)
```

```
Out[ ]: <Axes: >
```



Feature Engineering

```
In [ ]: # Extract year, month, day, and day of the week
sickness_data['year'] = sickness_data['date'].dt.year
sickness_data['month'] = sickness_data['date'].dt.month
sickness_data['day'] = sickness_data['date'].dt.day
sickness_data['day_of_week'] = sickness_data['date'].dt.dayofweek

# Create a binary feature indicating if the day is a weekend
sickness_data['is_weekend'] = sickness_data['day_of_week'].apply(lambda x: 1 if x > 6 else 0)

# Extract quarter
sickness_data['quarter'] = sickness_data['date'].dt.quarter

In [ ]: # Create lagged features for 'n_sick' and 'calls' for 1 day and 2 days
sickness_data['n_sick_lag1'] = sickness_data['n_sick'].shift(1)
sickness_data['n_sick_lag2'] = sickness_data['n_sick'].shift(2)
sickness_data['calls_lag1'] = sickness_data['calls'].shift(1)
sickness_data['calls_lag2'] = sickness_data['calls'].shift(2)

In [ ]: # Create rolling mean and standard deviation for 'n_sick' and 'calls' over a 7-day
sickness_data['n_sick_roll_mean'] = sickness_data['n_sick'].rolling(window=7).mean()
sickness_data['n_sick_roll_std'] = sickness_data['n_sick'].rolling(window=7).std()
sickness_data['calls_roll_mean'] = sickness_data['calls'].rolling(window=7).mean()
sickness_data['calls_roll_std'] = sickness_data['calls'].rolling(window=7).std()

In [ ]: # Calculate day-to-day difference for 'n_sick' and 'calls'
sickness_data['n_sick_diff'] = sickness_data['n_sick'].diff()
```



```
sickness_data['calls_diff'] = sickness_data['calls'].diff()
```

```
In [ ]: # Sick to Available Ratio: Ratio of drivers who called in sick to the total number
sickness_data['sick_to_available_ratio'] = sickness_data['n_sick'] / (sickness_data

# Emergency Call to Driver Ratio: Ratio of emergency calls to the total number of c
sickness_data['calls_to_driver_ratio'] = sickness_data['calls'] / (sickness_data['r
```

```
In [ ]: # Interaction between the number of drivers who called in sick and the number of en
sickness_data['sick_calls_interaction'] = sickness_data['n_sick'] * sickness_data['

# Interaction between the number of emergency calls and available drivers (both on
sickness_data['calls_driver_interaction'] = sickness_data['calls'] * (sickness_data
```

```
In [ ]: pd.set_option('display.max_columns', 30)

sickness_data.head(10)
```

```
Out[ ]:
```

	date	n_sick	calls	n_duty	n_sby	sby_need	dafted	year	month	day	day_of_week	is_we
0	2016-04-01	73	8154.0	1700	90	4.0	0.0	2016	4	1	4	
1	2016-04-02	64	8526.0	1700	90	70.0	0.0	2016	4	2	5	
2	2016-04-03	68	8088.0	1700	90	0.0	0.0	2016	4	3	6	
3	2016-04-04	71	7044.0	1700	90	0.0	0.0	2016	4	4	0	
4	2016-04-05	63	7236.0	1700	90	0.0	0.0	2016	4	5	1	
5	2016-04-06	70	6492.0	1700	90	0.0	0.0	2016	4	6	2	
6	2016-04-07	64	6204.0	1700	90	0.0	0.0	2016	4	7	3	
7	2016-04-08	62	7614.0	1700	90	0.0	0.0	2016	4	8	4	
8	2016-04-09	51	5706.0	1700	90	0.0	0.0	2016	4	9	5	
9	2016-04-10	54	6606.0	1700	90	0.0	0.0	2016	4	10	6	

```
In [ ]: sickness_data.shape
```

```
Out[ ]: (1152, 27)
```

Handling Missing Values

```
In [ ]: # Columns for which to apply forward fill and then backward fill
columns_to_ffill = [
    'n_sick_lag1', 'n_sick_lag2', 'calls_lag1', 'calls_lag2',
    'n_sick_diff', 'calls_diff'
]
```

```

# Apply both forward fill and backward fill for these columns
for column in columns_to_ffill:
    sickness_data[column] = sickness_data[column].ffill().bfill()

# Columns for which to apply backward fill
rolling_columns_to_bfill = [
    'n_sick_roll_mean', 'n_sick_roll_std', 'calls_roll_mean', 'calls_roll_std'
]

# Apply backward fill for these columns
for column in rolling_columns_to_bfill:
    sickness_data[column] = sickness_data[column].bfill()

# (Optional) Verify there are no more NaN values
nan_after_handling = sickness_data.isna().sum()
print(nan_after_handling)

```

```

date                0
n_sick              0
calls               0
n_duty              0
n_sby               0
sby_need            0
dafted              0
year                0
month               0
day                 0
day_of_week         0
is_weekend          0
quarter             0
n_sick_lag1         0
n_sick_lag2         0
calls_lag1          0
calls_lag2          0
n_sick_roll_mean    0
n_sick_roll_std     0
calls_roll_mean     0
calls_roll_std      0
n_sick_diff         0
calls_diff          0
sick_to_available_ratio 0
calls_to_driver_ratio 0
sick_calls_interaction 0
calls_driver_interaction 0
dtype: int64

```

```

In [ ]: # 1. Temporal Visualization of New Features:
plt.figure(figsize=(14, 6))
plt.plot(sickness_data['date'], sickness_data['n_sick_roll_mean'], label='Rolling Mean')
plt.plot(sickness_data['date'], sickness_data['n_sick_diff'], label='Difference of')
plt.title('Rolling Mean & Difference of Drivers Sick')
plt.legend()
plt.show()

plt.figure(figsize=(14, 6))
plt.plot(sickness_data['date'], sickness_data['calls_roll_mean'], label='Rolling Mean')
plt.plot(sickness_data['date'], sickness_data['calls_diff'], label='Difference of')
plt.title('Rolling Mean & Difference of Emergency Calls')
plt.legend()
plt.show()

# 2. Correlation Analysis:

```

```

correlations = sickness_data.drop(columns=['date']).corr()['sby_need'].sort_values()
plt.figure(figsize=(12, 6))
correlations.plot(kind='bar', color='coral')
plt.title('Correlation with Number of Standbys Activated')
plt.show()

# 3. Distribution Analysis:
plt.figure(figsize=(8, 6))
sns.histplot(sickness_data['n_sick_roll_mean'], kde=True, color='skyblue', bins=30)
plt.title('Distribution of Rolling Mean of Drivers Sick')
plt.show()

# 4. Seasonal Patterns:
plt.figure(figsize=(8, 6))
month_avg_sick = sickness_data.groupby('month')['n_sick'].mean()
month_avg_sick.plot(kind='bar', color='lightgreen')
plt.title('Average Number of Drivers Sick by Month')
plt.xticks(ticks=range(12), labels=month_avg_sick.index, rotation=0)
plt.show()


# 5. Weekday vs. Weekend Analysis:
plt.figure(figsize=(8, 6))
weekday_avg_sick = sickness_data.groupby('day_of_week')['n_sick'].mean()
weekday_avg_sick.plot(kind='bar', color='lightblue')
plt.title('Average Number of Drivers Sick by Day of Week')
plt.xticks(ticks=range(7), labels=['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun'])
plt.show()

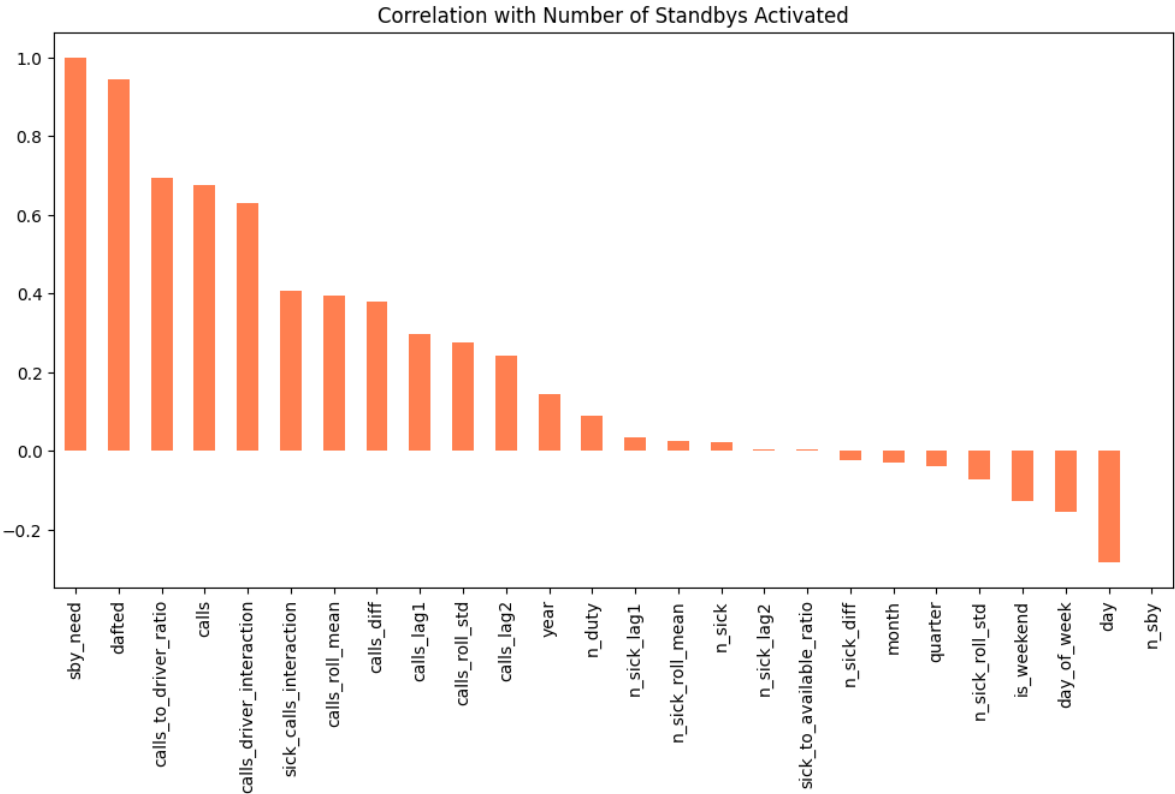

# 6. Box Plots:
plt.figure(figsize=(12, 6))
plt.subplot(1, 2, 1)
sns.boxplot(data=sickness_data, y='n_sick_roll_mean', color='pink')
plt.title('Box Plot of Rolling Mean of Drivers Sick')



plt.subplot(1, 2, 2)
sns.boxplot(data=sickness_data, y='calls_diff', color='yellow')
plt.title('Box Plot of Difference of Emergency Calls')
plt.tight_layout()
plt.show()

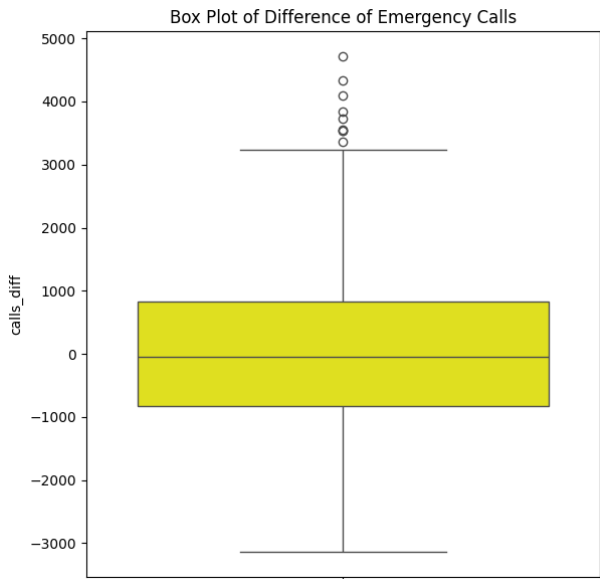
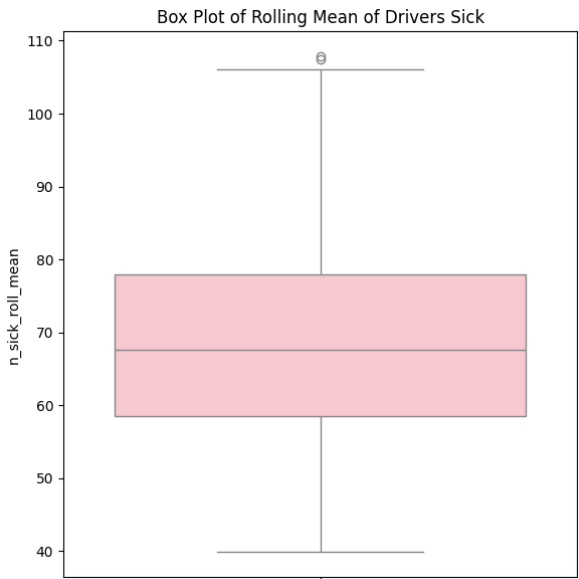

# 7. Heatmap of Correlations:
plt.figure(figsize=(12, 10))
correlation_matrix = sickness_data.drop(columns=['date']).corr()
sns.heatmap(correlation_matrix, cmap='coolwarm', annot=True, fmt=".2f")
plt.title('Heatmap of Feature Correlations')
plt.show()


```





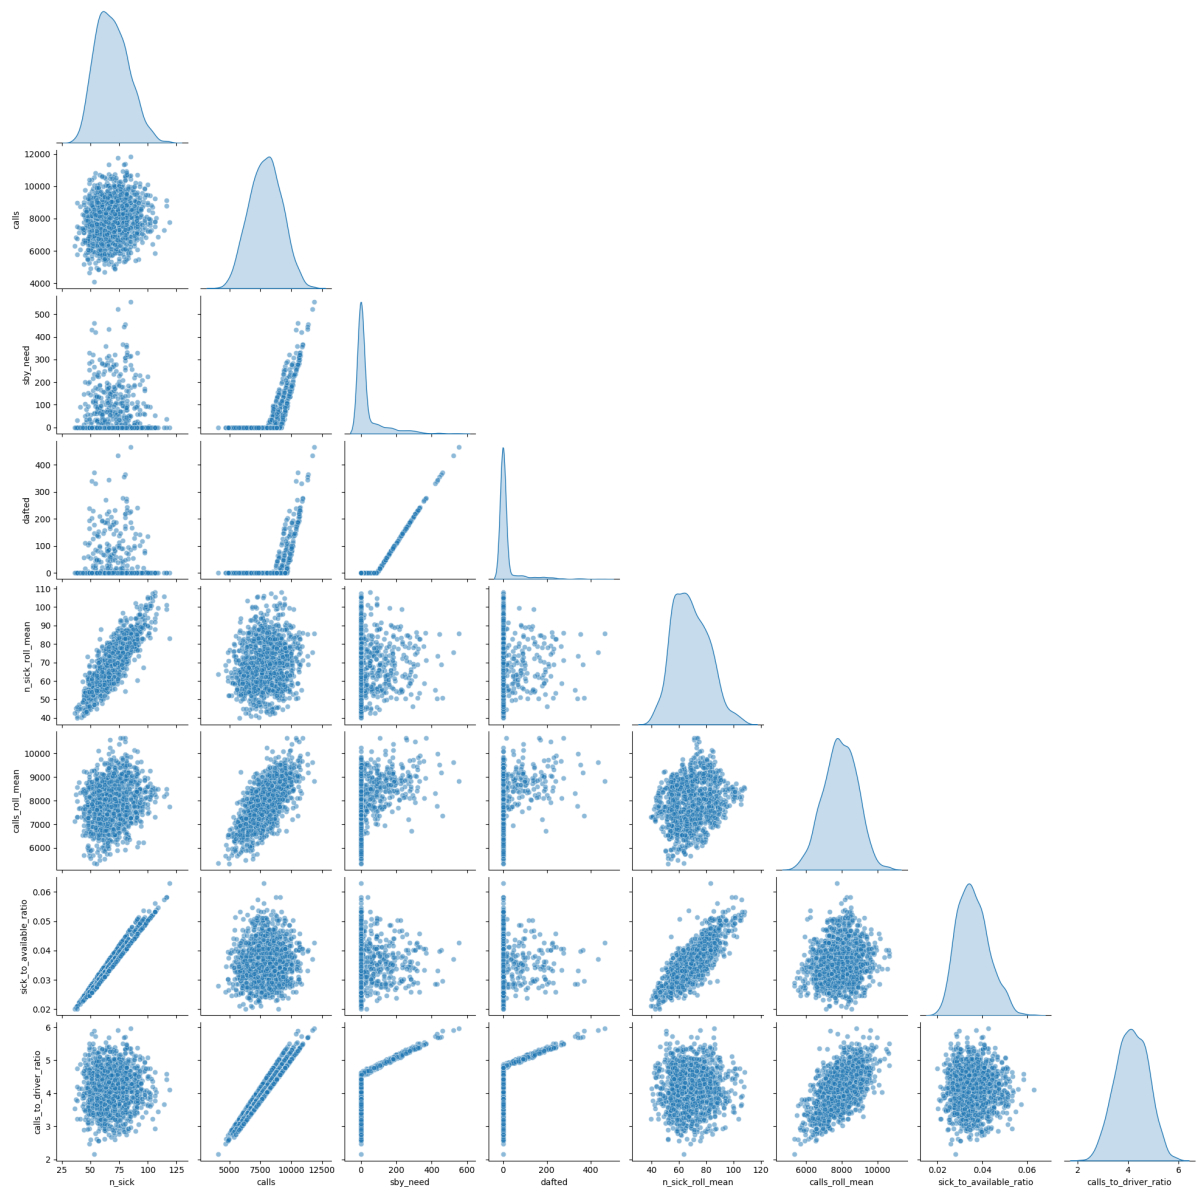






```
In [ ]: # Select relevant columns for pair plot
selected_columns = ['n_sick', 'calls', 'sby_need', 'dafted', 'n_sick_rol_mean',
                    'calls_rol_mean', 'sick_to_available_ratio', 'calls_to_driver_ratio']

# Generate pair plot for selected columns
pairplot_data = sickness_data[selected_columns]
sns.pairplot(pairplot_data, corner=True, diag_kind='kde', markers='o', plot_kws={'alpha': 0.5})
plt.suptitle('Pair Plot for Selected Features', y=1.02)
plt.show()
```



Linear Regression Model

```
In [ ]: # Define features and target variable
X = sickness_data.drop(columns=['date', 'sby_need'])
y = sickness_data['sby_need']

# Split the data into training and testing sets (80% training, 20% testing)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Initialize and train the baseline model (Linear Regression)
baseline_model = LinearRegression()
baseline_model.fit(X_train, y_train)

# Predict on the training set
y_pred_train = baseline_model.predict(X_train)

# Evaluate the baseline model's performance on training data
mae_train = mean_absolute_error(y_train, y_pred_train)
mse_train = mean_squared_error(y_train, y_pred_train)
rmse_train = mean_squared_error(y_train, y_pred_train, squared=False)
r2_train = r2_score(y_train, y_pred_train)

# Predict on the test set
```



```

y_pred = baseline_model.predict(X_test)

# Evaluate the baseline model
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = mean_squared_error(y_test, y_pred, squared=False)
r2 = r2_score(y_test, y_pred)

print("Performance on Training Data:")
print(f"Mean Absolute Error (MAE): {mae_train}")
print(f"Mean Squared Error (MSE): {mse_train}")
print(f"Root Mean Squared Error (RMSE): {rmse_train}")
print(f"R^2 (Coefficient of Determination): {r2_train}")
print("\nPerformance on Test Data:")
print(f"Mean Absolute Error (MAE): {mae}")
print(f"Mean Squared Error (MSE): {mse}")
print(f"Root Mean Squared Error (RMSE): {rmse}")
print(f"R^2 (Coefficient of Determination): {r2}")

```

Performance on Training Data:
Mean Absolute Error (MAE): 16.86249271405271
Mean Squared Error (MSE): 445.7586201349591
Root Mean Squared Error (RMSE): 21.11299647456417
R^2 (Coefficient of Determination): 0.9359094318861493

Performance on Test Data:
Mean Absolute Error (MAE): 15.959891029202183
Mean Squared Error (MSE): 397.75521259023304
Root Mean Squared Error (RMSE): 19.943801357570553
R^2 (Coefficient of Determination): 0.8949557636782547

Gradient Boosting Machines

```

In [ ]: # Initialize the GBM model
gbm_model = GradientBoostingRegressor(n_estimators=100, learning_rate=0.1, max_dept

# Train the GBM model on the training data
gbm_model.fit(X_train, y_train)

# Predict on the training set
y_pred_train_gbm = gbm_model.predict(X_train)

# Evaluate the GBM model's performance on training data
mae_gbm_train = mean_absolute_error(y_train, y_pred_train_gbm)
mse_gbm_train = mean_squared_error(y_train, y_pred_train_gbm)
rmse_gbm_train = mean_squared_error(y_train, y_pred_train_gbm, squared=False)
r2_gbm_train = r2_score(y_train, y_pred_train_gbm)

# Predict on the test set
y_pred_gbm = gbm_model.predict(X_test)

# Evaluate the GBM model's performance
mae_gbm = mean_absolute_error(y_test, y_pred_gbm)
mse_gbm = mean_squared_error(y_test, y_pred_gbm)
rmse_gbm = mean_squared_error(y_test, y_pred_gbm, squared=False)
r2_gbm = r2_score(y_test, y_pred_gbm)

print("Performance on Training Data:")
print(f"Mean Absolute Error (MAE): {mae_gbm_train}")
print(f"Mean Squared Error (MSE): {mse_gbm_train}")
print(f"Root Mean Squared Error (RMSE): {rmse_gbm_train}")

```

```
print(f"R^2 (Coefficient of Determination): {r2_gbm_train}")
print("\nPerformance on Test Data:")
print(f"Mean Absolute Error (MAE): {mae_gbm}")
print(f"Mean Squared Error (MSE): {mse_gbm}")
print(f"Root Mean Squared Error (RMSE): {rmse_gbm}")
print(f"R^2 (Coefficient of Determination): {r2_gbm}")
```

Performance on Training Data:

Mean Absolute Error (MAE): 0.4979518677197901

Mean Squared Error (MSE): 1.2840423087438884

Root Mean Squared Error (RMSE): 1.1331559066359265

R^2 (Coefficient of Determination): 0.9998153821433118

Performance on Test Data:

Mean Absolute Error (MAE): 1.1516967679975911

Mean Squared Error (MSE): 11.82975794674722

Root Mean Squared Error (RMSE): 3.4394415166923857

R^2 (Coefficient of Determination): 0.9968758476317762

```
In [ ]: from sklearn.model_selection import RandomizedSearchCV
from sklearn.ensemble import GradientBoostingRegressor

# Simplified and adjusted hyperparameters grid
simplified_param_grid = {
    'n_estimators': [50, 100],
    'learning_rate': [0.01, 0.1],
    'max_depth': [2, 3],
    'subsample': [0.8, 1.0],
    'max_features': ['sqrt', 'log2', None]
}

# Initialize the RandomizedSearchCV object with the simplified hyperparameters
simplified_random_search = RandomizedSearchCV(GradientBoostingRegressor(random_state=42,
    param_distributions=simplified_param_grid,
    n_iter=10, # Reduced number of iterations
    scoring='neg_mean_squared_error',
    n_jobs=-1,
    cv=2, # Reduced CV folds
    random_state=42,
    error_score='raise')

# Fit to the training data
simplified_random_search.fit(X_train, y_train)

# Extract the best estimator after the search
best_gbm_simplified = simplified_random_search.best_estimator_
print(best_gbm_simplified)

GradientBoostingRegressor(max_depth=2, max_features='sqrt', random_state=42,
    subsample=0.8)
```

```
In [ ]: # Predict on the training data using the best GBM model
y_train_pred = best_gbm_simplified.predict(X_train)

# Evaluate the model's performance on training data
mae_train = mean_absolute_error(y_train, y_train_pred)
mse_train = mean_squared_error(y_train, y_train_pred)
rmse_train = mean_squared_error(y_train, y_train_pred, squared=False)
r2_train = r2_score(y_train, y_train_pred)

# Predict on the test data using the best GBM model
y_test_pred = best_gbm_simplified.predict(X_test)

# Evaluate the model's performance on test data
mae_test = mean_absolute_error(y_test, y_test_pred)
```

```

mse_test = mean_squared_error(y_test, y_test_pred)
rmse_test = mean_squared_error(y_test, y_test_pred, squared=False)
r2_test = r2_score(y_test, y_test_pred)

print("Performance on Training Data:")
print(f"Mean Absolute Error (MAE): {mae_train}")
print(f"Mean Squared Error (MSE): {mse_train}")
print(f"Root Mean Squared Error (RMSE): {rmse_train}")
print(f"R^2 (Coefficient of Determination): {r2_train}")
print("\nPerformance on Test Data:")
print(f"Mean Absolute Error (MAE): {mae_test}")
print(f"Mean Squared Error (MSE): {mse_test}")
print(f"Root Mean Squared Error (RMSE): {rmse_test}")
print(f"R^2 (Coefficient of Determination): {r2_test}")

```

Performance on Training Data:
Mean Absolute Error (MAE): 4.306941597889101
Mean Squared Error (MSE): 41.7662257175321
Root Mean Squared Error (RMSE): 6.462679453410335
R^2 (Coefficient of Determination): 0.9939949088737813

Performance on Test Data:
Mean Absolute Error (MAE): 5.613848985003764
Mean Squared Error (MSE): 119.29142931425552
Root Mean Squared Error (RMSE): 10.922061587184698
R^2 (Coefficient of Determination): 0.9684960078575897

```

In [ ]: # Predicted values for Linear Regression
y_pred_lr = baseline_model.predict(X_test)

# Predicted values for untuned GBM
y_pred_gbm = gbm_model.predict(X_test)

# Predicted values for hyperparameter-tuned GBM
y_pred_best_gbm = best_gbm_simplified.predict(X_test)

# Function to check insufficient standby predictions
def check_insufficient_standbys(predictions, actual_sby_need, n_sby_values):
    insufficient_standbys = sum(predictions > n_sby_values)
    actual_insufficient_standbys = sum(actual_sby_need > n_sby_values)
    return insufficient_standbys, actual_insufficient_standbys

# Check for each model
insufficient_lr, actual_insufficient = check_insufficient_standbys(y_pred_lr, y_test, X_test['n_sby'])
insufficient_gbm, _ = check_insufficient_standbys(y_pred_gbm, y_test, X_test['n_sby'])
insufficient_best_gbm, _ = check_insufficient_standbys(y_pred_best_gbm, y_test, X_test['n_sby'])

print(f"Days with insufficient standbys (Actual): {actual_insufficient}")
print(f"Days with insufficient standbys (Linear Regression): {insufficient_lr}")
print(f"Days with insufficient standbys (Untuned GBM): {insufficient_gbm}")
print(f"Days with insufficient standbys (Tuned GBM): {insufficient_best_gbm}")

```

Days with insufficient standbys (Actual): 25
Days with insufficient standbys (Linear Regression): 18
Days with insufficient standbys (Untuned GBM): 25
Days with insufficient standbys (Tuned GBM): 25

```

In [ ]: import matplotlib.pyplot as plt

# Compute feature importance
# For Linear Regression: Coefficients as feature importance
lr_importance = baseline_model.coef_

# For Gradient Boosting Machines: Feature importance attribute
gbm_importance = gbm_model.feature_importances_


```

```
# Visualization
fig, ax = plt.subplots(2, 1, figsize=(15, 12))

# Plotting feature importance for Linear Regression
ax[0].bar(X.columns, lr_importance, color='skyblue')
ax[0].set_title('Feature Importance (Linear Regression)')
ax[0].tick_params(axis='x', rotation=45)
ax[0].set_ylabel('Coefficient Value')

# Plotting feature importance for GBM
ax[1].bar(X.columns, gbm_importance, color='salmon')
ax[1].set_title('Feature Importance (GBM)')
ax[1].tick_params(axis='x', rotation=45)
ax[1].set_ylabel('Importance Value')

plt.tight_layout()
plt.show()
```




Error Analysis

```
In [ ]: import numpy as np

# Calculate residuals for the GBM model
residuals = y_test - y_test_pred

# Plot residuals
plt.figure(figsize=(10, 6))
plt.scatter(y_test_pred, residuals, alpha=0.5, color='darkred')
plt.axhline(y=0, color='black', linestyle='--')
plt.title('Residuals vs. Predicted Values')
plt.xlabel('Predicted Values')
```

```
plt.ylabel('Residuals')
plt.show()
```



```
In [ ]: # Calculate residuals for the GBM model
residuals = y_test - y_test_pred


# Determine the threshold for high errors
threshold = 2 * np.std(residuals)
high_error_indices = np.where(np.abs(residuals) > threshold)[0]

# Extract high error data from the test set
high_error_data = X_test.iloc[high_error_indices]

# Visualize the distribution of residuals and highlight the high-error instances
plt.figure(figsize=(14, 6))
plt.hist(residuals, bins=50, color='lightblue', label='All Residuals', alpha=0.7)
plt.hist(residuals.iloc[high_error_indices], bins=50, color='darkred', label='High-')
plt.axvline(x=threshold, color='grey', linestyle='--', label=f'+2 Std Dev ({threshold}')
plt.axvline(x=-threshold, color='grey', linestyle='--', label=f'-2 Std Dev (-{threshold}')
plt.title('Distribution of Residuals')
plt.xlabel('Residual Value')
plt.ylabel('Frequency')
plt.legend()
plt.show()

# To inspect the high-error data, you can further examine the `high_error_data` Data
# For a more detailed view, add the residuals to this data:
high_error_data['Residual'] = residuals.iloc[high_error_indices].values

# Display the first few rows of the high-error data for inspection
print(high_error_data.head())
```



	n_sick	calls	n_duty	n_sby	dafted	year	month	day	day_of_week	\
174	77	8286.0	1700	90	0.0	2016	9	22	3	
101	70	9492.0	1700	90	179.0	2016	7	11	0	
1006	82	9672.0	1900	90	27.0	2019	1	2	2	
1099	67	9444.0	1900	90	0.0	2019	4	5	4	
96	51	9702.0	1700	90	202.0	2016	7	6	2	

	is_weekend	quarter	n_sick_lag1	n_sick_lag2	calls_lag1	calls_lag2	\
174	0	3	80.0	72.0	8532.0	7362.0	
101	0	3	70.0	69.0	7374.0	8532.0	
1006	0	1	57.0	75.0	8382.0	6318.0	
1099	0	2	75.0	79.0	10260.0	11328.0	
96	0	3	47.0	50.0	8550.0	8526.0	

	n_sick_roll_mean	n_sick_roll_std	calls_roll_mean	calls_roll_std	\
174	75.714286	3.817254	7530.857143	873.224566	
101	61.571429	9.396048	8730.857143	775.149755	
1006	74.285714	8.056349	7638.000000	1465.646615	
1099	80.000000	7.416198	9756.000000	1879.590381	
96	46.142857	4.220133	8245.714286	1016.203017	

	n_sick_diff	calls_diff	sick_to_available_ratio	calls_to_driver_ratio	\
174	-3.0	-246.0	0.043017	4.629050	
101	0.0	2118.0	0.039106	5.302793	
1006	25.0	1290.0	0.041206	4.860302	
1099	-8.0	-816.0	0.033668	4.745729	
96	4.0	1152.0	0.028492	5.420112	

	sick_calls_interaction	calls_driver_interaction	Residual
174	638022.0	14831940.0	23.079963
101	664440.0	16990680.0	44.238582
1006	793104.0	19247280.0	-64.973504
1099	632748.0	18793560.0	-35.971229
96	494802.0	17366580.0	68.091800

C:\Users\s9\AppData\Local\Temp\ipykernel_6408\3816389916.py:25: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using `.loc[row_indexer,col_indexer] = value` instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

`high_error_data['Residual'] = residuals.iloc[high_error_indices].values`

In []: